

# ARTIFICIAL INTELLIGENCE – FAKE NEWS DETECTION USING NLP

## PHASE 5

### PROBLEM STATEMENT:

The problem is to develop a fake news detection model using a Kaggle dataset. The goal is to distinguish between genuine and fake news articles based on their titles and text. This project involves using natural language processing (NLP) techniques to pre-process the text data, building a machine learning model for classification, and evaluating the model's performance.

### DESIGN THINKING PROCESS:

#### 1. DATA SOURCE

Action: Choose the Kaggle dataset containing news articles, titles, and labels (genuine or fake).

Rationale: Selecting a relevant and high-quality dataset is the foundation of building an effective fake news detection model. Kaggle provides a diverse range of datasets suitable for this task.

#### 2. DATA PREPROCESSING

Action: Clean and preprocess the textual data to prepare it for analysis.

Rationale: Data preprocessing is essential to ensure that the text data is consistent, free from noise, and suitable for further analysis. Steps may include:

- Removing special characters, punctuation, and HTML tags.
- Tokenization (splitting text into words or tokens).
- Lowercasing all text to ensure consistency.
- Removing stop words (common words like 'the', 'and', 'is', etc. that don't carry significant meaning).
- Handling missing data if any.

#### 3. FEATURE EXTRACTION

Action: Utilize techniques like TF-IDF (Term Frequency-Inverse Document Frequency) or word embeddings to convert text into numerical features.

Rationale: Converting text into numerical features is necessary for machine learning models to process the data. TF-IDF and word embeddings like Word2Vec or GloVe can capture semantic information and relationships between words, making them suitable for this task.

#### **4. MODEL SELECTION**

Action: Select a suitable classification algorithm (e.g., Logistic Regression, Random Forest, or Neural Networks) for the fake news detection task.

Rationale: The choice of a classification algorithm can significantly impact the model's performance. Different algorithms have different strengths and weaknesses, and the choice should be based on the dataset's characteristics and complexity. Commonly used algorithms for text classification include Logistic Regression, Random Forest, and Neural Networks. Evaluating multiple algorithms may be necessary to identify the best-performing one.

#### **5. MODEL TRAINING**

Action: Train the selected model using the preprocessed data.

Rationale: Model training involves feeding the preprocessed data into the chosen classification algorithm. The model learns to identify patterns and relationships between the text features and their corresponding labels (genuine or fake news).

#### **6. EVALUATION**

Action: Evaluate the model's performance using metrics like accuracy, precision, recall, F1-score, and ROC-AUC.

Rationale: Model evaluation is crucial to assess how well the fake news detection system is performing. The selected metrics provide insights into various aspects of the model's performance:

- Accuracy: Overall correctness of the model's predictions.
- Precision: Proportion of true positive predictions among all positive predictions, measuring the model's ability to avoid false positives.

- Recall: Proportion of true positive predictions among all actual positives, measuring the model's ability to capture genuine fake news.
- F1-score: Harmonic mean of precision and recall, offering a balanced measure of model performance.
- ROC-AUC: Receiver Operating Characteristic - Area Under the Curve measures the model's ability to distinguish between the two classes.

## **ENHANCEMENTS USING BERT AND LSTM**

To further enhance the accuracy of the fake news detection model, consider incorporating advanced techniques such as BERT (Bidirectional Encoder Representations from Transformers) and LSTM (Long Short-Term Memory) networks.

### **7. BERT INTEGRATION**

Action: Implement BERT-based models for feature extraction.

Rationale: BERT, as a pre-trained transformer model, captures contextual information and relationships between words in a sentence. By using BERT embeddings, the model can better understand the semantics of the text, improving the representation of complex language constructs in news articles.

### **8. LSTM INTEGRATION**

Action: Incorporate LSTM layers into the neural network architecture.

Rationale: LSTM networks are effective in capturing sequential dependencies in data. Since news articles often have a sequential structure, LSTM layers can enhance the model's ability to understand the temporal aspects of language. This is particularly useful for detecting nuances and context in fake news articles.

### **9. MODEL TRAINING AND EVALUATION**

Action: Retrain the model using the enhanced feature extraction techniques (BERT and LSTM).

Rationale: By leveraging BERT and LSTM, the model can potentially achieve higher accuracy and better generalization to complex patterns in the data. Retrain the model and evaluate its performance using the defined metrics to ensure improvements in fake news detection accuracy.

## DATASET LINK:

<https://www.kaggle.com/clmentbisailon/fake-and-real-news-dataset>

## DESCRIPTION:

The significance of social media has increased manifold in the past few decades as it helps people from even the most remote corners of the world stay connected. With the COVID-19 pandemic raging, social media has become more relevant and widely used than ever before, and along with this, there has been a resurgence in the circulation of fake news and tweets that demand immediate attention. In this paper, we describe our Fake News Detection system that automatically identifies whether a tweet related to COVID-19 is "real" or "fake", as a part of CONSTRAINT COVID19 Fake News Detection in English challenge. We have used an ensemble model consisting of pre-trained models that has helped us achieve a joint 8th position on the leader board. We have achieved an F1-score of 0.9831 against a top score of 0.9869. Post completion of the competition, we have been able to drastically improve our system by incorporating a novel heuristic algorithm based on username handles and link domains in tweets fetching an F1-score of 0.9883 and achieving state-of-the art results on the given dataset.

## FAKE NEWS DETECTION USING NLP

### Fake News Detection

**DESCRIPTION:**The problem at hand is to develop a fake news detection model using a dataset obtained from Kaggle. The objective is to create a system that can effectively distinguish between genuine and fake news articles based on their titles and textual content. This project requires the utilization of Natural Language Processing (NLP) techniques to preprocess and transform the text data, building a machine learning model for classification, and subsequently evaluating the model's performance

Dataset Link: <https://www.kaggle.com/datasets/clmentbisailon/fake-and-real-news-dataset>

### DATA LOADING

```
import pandas as pd

# Load the dataset (replace 'path_to_fake_news_dataset.csv' with the
# actual path to your dataset)
dataset_path = 'path_to_fake_news_dataset.csv'
fake_data = pd.read_csv(dataset_path)

# Display the first few rows of the dataset to inspect the data
fake_data.head()
```

Here, we need to replace 'path\_to\_fake\_news\_dataset.csv' with the actual file path where we stored the dataset. This code will load the dataset into a Pandas DataFrame and display the first few rows to inspect the data.

### TEXT PROCESSING

```
import nltk
from sklearn.preprocessing import LabelEncoder
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize

# Download NLTK data (if not already downloaded)
nltk.download('stopwords')
nltk.download('punkt')

# Remove missing values (if any)
fake_data = fake_data.dropna()

# Combine 'title' and 'text' columns for analysis
fake_data['text'] = fake_data['title'] + " " + fake_data['text']

# Tokenization and stop-word removal
stop_words = set(stopwords.words('english'))
fake_data['text'] = fake_data['text'].apply(lambda x: ' '.join([word
for word in word_tokenize(x) if word.lower() not in stop_words]))
```

```
# Label encoding (fake: 1, real: 0)
label_encoder = LabelEncoder()
fake_data['label'] = label_encoder.fit_transform(fake_data['label'])

# Display the preprocessed dataset
fake_data.head()
```

Removes any missing values in the dataset. Combines the 'title' and 'text' columns into a single 'text' column for analysis. Tokenizes the text using NLTK's word\_tokenize function. Removes English stopwords from the tokenized text. Performs label encoding to convert 'fake' and 'real' labels to numeric values ('fake' as 1 and 'real' as 0).

## FEATURE EXTRACTION

```
from sklearn.feature_extraction.text import TfidfVectorizer

# Create a TF-IDF vectorizer with a maximum of 5000 features
tfidf_vectorizer = TfidfVectorizer(max_features=5000)

# Apply TF-IDF vectorization to the 'text' column
X = tfidf_vectorizer.fit_transform(fake_data['text'])

# Display the TF-IDF features
X.toarray() # Converts the sparse matrix to a dense array for display
            (use carefully)

# Display the feature names
feature_names = tfidf_vectorizer.get_feature_names_out()
print("Example Feature Names:", feature_names[:10])
```

This code utilizes the TfidfVectorizer from scikit-learn to convert the text data into a TF-IDF feature matrix. The max\_features parameter limits the number of features to 5000, but you can adjust this number as needed.

The resulting X contains the TF-IDF features for each document in your dataset. The feature names can be obtained using get\_feature\_names\_out(). Please note that displaying the entire feature matrix (X.toarray()) may not be practical for a large dataset, so use it judiciously.

## MODEL TRAINING

```
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X,
fake_data['label'], test_size=0.2, random_state=42)

# Create a Multinomial Naive Bayes classifier
classifier = MultinomialNB()
```

```
# Train the classifier on the training data
classifier.fit(X_train, y_train)
```

In this code, we use the `train_test_split` function to split the TF-IDF features (X) and the labels (`fake_data['label']`) into training and testing sets. We then create a Multinomial Naive Bayes classifier (`MultinomialNB`) and train it using the training data.

## MODEL EVALUATION

```
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# Make predictions on the test set
y_pred = classifier.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)

# Generate confusion matrix
confusion_mat = confusion_matrix(y_test, y_pred)

# Generate classification report
classification_rep = classification_report(y_test, y_pred)

# Print evaluation results
print(f"Accuracy: {accuracy}")
print("Confusion Matrix:")
print(confusion_mat)
print("Classification Report:")
print(classification_rep)

# Plot a confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_mat, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Real', 'Fake'], yticklabels=['Real', 'Fake'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

In this code, we first make predictions on the test set using the trained Multinomial Naive Bayes classifier. Then, we calculate the accuracy of the model, generate a confusion matrix, and create a classification report. Finally, we print the results and plot the confusion matrix for visualization.

**CONCLUSION:** The problem is to develop a fake news detection model using a Kaggle dataset. The goal is to distinguish between genuine and fake news articles based on their titles and text. This project involves using natural

language processing (NLP) techniques to pre-process the text data, building a machine learning model for classification, and evaluating the model's performance.